

Self-Study

1. Why must operator overloading be static and non-private ?

In C#, operator overloading must be static and non-private for several key reasons:

1. **Static Requirement:**

Operators work on operands, which may be instances or types, but the operator itself is tied to the type rather than an instance.

Defining operator overloads as static methods ensures that the operator applies uniformly to the provided operands without relying on instance state. This also aligns with how operators conceptually work as binary or unary operations performed on values, not on specific object instances. Static methods can be called without an instance, matching the expectation that operators are used like language-level operations.

2. **Non-Private Accessibility:**

Operator overloads must be accessible publicly (or at least non-private) so that the overloaded operator can be used anywhere the type is visible. Operators are syntactic sugar around method calls; if the operator method were private, the compiler and external code could not invoke it naturally using operator syntax, breaking usability and expected behavior.

Together, these constraints ensure that the operator behaves like a natural extendable part of the language syntax, is consistently applicable at the type level, and is accessible where needed. This provides type safety, compile-time resolution, and clear semantics in code using overloaded operators in C#.

2. What is the backing field and how does it work in intermediate language?

A backing field, also known as a **private member variable**, is a private variable that stores the value of a property. It's used when a property's `get` or `set` accessor contains logic beyond simply returning or setting a value.

How It Relates to the Intermediate Language

The backing field's relationship to the intermediate language (IL) is fundamental to how .NET works. When you create an **auto-implemented property** in C# (e.g., `public int MyProperty { get; set; }`), the compiler automatically generates a private, anonymous backing field in the IL.

This means that even though you don't explicitly declare the backing field in your C# code, it's there at the IL level. The `get` accessor of the property is compiled into IL that loads the value from this hidden backing field, while the `set` accessor is compiled into IL that stores the new value into it.

If you write a property with a **manual backing field** (e.g., `private int myField; public int MyProperty { get { return myField; } set { myField = value; } }`), the compiler uses your explicitly declared field. The generated IL for the property's accessors directly references this field.

In both cases, the backing field is the storage location for the property's value, and its management is a key part of the .NET compiler's work in translating high-level code into low-level IL instructions. The concept of a backing field is an essential part of the IL code that makes properties function.

3. What is Record in C#9.0 ?

In C# 9.0, a record is a new reference type designed for encapsulating immutable data more easily and concisely compared to traditional classes. Records provide built-in functionality focused on data-centric scenarios with features such as:

- **Value-based equality:** Two record instances are considered equal if all their property values are equal, unlike classes which use reference equality by default.
- **Immutable properties:** With positional syntax or init-only properties, records encourage immutability, making data safer and less error-prone.
- **Concise syntax:** Records support a compact declaration style, often with positional parameters that automatically generate properties.
- **Built-in formatting:** Compiler generates a ToString() override that displays property names and values for easy debugging and logging.
- **Non-destructive mutation:** The with expression creates new copies of records with modified properties without changing the original instance.
- **Support for inheritance:** Records can inherit from other records, allowing hierarchical data models with value-based equality and immutability.

Records are primarily intended for scenarios such as data transfer objects (DTOs), configuration data, and other cases where immutable data with value semantics is beneficial, simplifying the creation and management of such data compared to traditional classes.

4. Why would it be better to start the enum with 1 in most use cases ?

It is often better to start an enum with 1 in most use cases for these reasons:

1. **Avoid Confusion with Default Value 0:** In C#, the default value for an uninitialized enum variable is 0. If 0 is not a meaningful or valid state for the enum, starting from 1 helps differentiate "no value" or "uninitialized" from valid enum values, preventing unintended behaviors.
2. **Intuition and Semantics:** Starting from 1 aligns more naturally with real-world counting and logical ordering where 1-based indexing is intuitive. For many enums, 0 may not correspond to any meaningful state, so starting from 1 makes enum values more semantically meaningful.
3. **Explicit Value Handling:** By starting at 1, any zero value can be reserved to represent "none," "unknown," or "uninitialized" explicitly, which can improve code clarity and validation.
4. **Avoid Problems in Flags Enums:** For enums representing bit flags, starting from 1 (which is 2^0) is crucial to ensure correct bit masking. Starting at 0 would cause the first flag to be 0, which is problematic logically.

However, in cases where 0 represents a valid default value or state, starting at 0 is appropriate. The choice depends on the semantics and usage context of the enum.

In summary, starting enums at 1 is generally better when zero does not represent a valid or useful state, helping avoid errors and making code logic clearer.

5. What are Enum Use Cases that need to manually assign values?

Use cases where enums require manually assigned values include:

1. **Mapping to External or Standardized Values:** When enum values correspond to external systems or standards like HTTP status codes (OK = 200, NotFound = 404), error codes, or protocol values, manual assignment ensures correctness and compatibility.
2. **Persisted or Serialized Data:** If enum values are stored or serialized, assigning fixed values prevents breaking data integrity if enum members are reordered or updated, avoiding mismatches during deserialization or reading from databases.
3. **Flags and Bitmask Operations:** Enums representing bit flags use powers of two (1, 2, 4, 8...) for bitwise operations; manual assignment of these values is necessary for correct flag combination and checks.
4. **Business Logic with Specific Numeric Meaning:** When numeric values carry semantic meaning or ranking, like priorities with explicit scores (Low = 1, Medium = 5, High = 10), manual assignments help convey and enforce these meanings.
5. **Backward Compatibility and Versioning:** When evolving an enum used across different application versions or distributed systems, fixed manual values ensure compatibility between components.
6. **Non-Sequential or Sparse Values:** For enums where values are non-sequential or sparsely distributed, manual assignment provides clarity and control.

In general, manual enum value assignment is favored when enumeration values must remain stable, interoperable, semantically precise, or optimized for particular operations like bitwise flags.

6. What can a static class contain rather than static members?

In C#, a static class can contain only static members. This means it cannot contain any instance members such as instance fields, properties, methods, or constructors. Specifically, a static class can contain:

- Static fields (static data members)
- Static methods
- Static properties
- Static events
- Static constructors (used to initialize static data)

Because a static class cannot be instantiated, allowing only static members ensures that all members are accessed at the class level rather than through an object instance. This is by language design to enforce that the class acts as a container for global or singleton-like functionality.

In summary, a static class cannot contain anything other than static members, no instance members are allowed.